



Introduction to Lustre


Mark Cheeseman
Mark.Cheeseman@csiro.au



Parallel filesystem basics

What is a filesystem?

- A way of organizing data on a storage device
- We see data as having two components:
 - **Data**: raw information
 - numbers, characters, images, etc
 - **Metadata**: description of the stored data
 - array length
 - type (integer, float, double, etc)
 - address at which data is stored
- I/O operations use both components



What's different with parallel filesystems?



A parallel filesystem employs both task and domain parallelism to increase throughput of I/O operations.

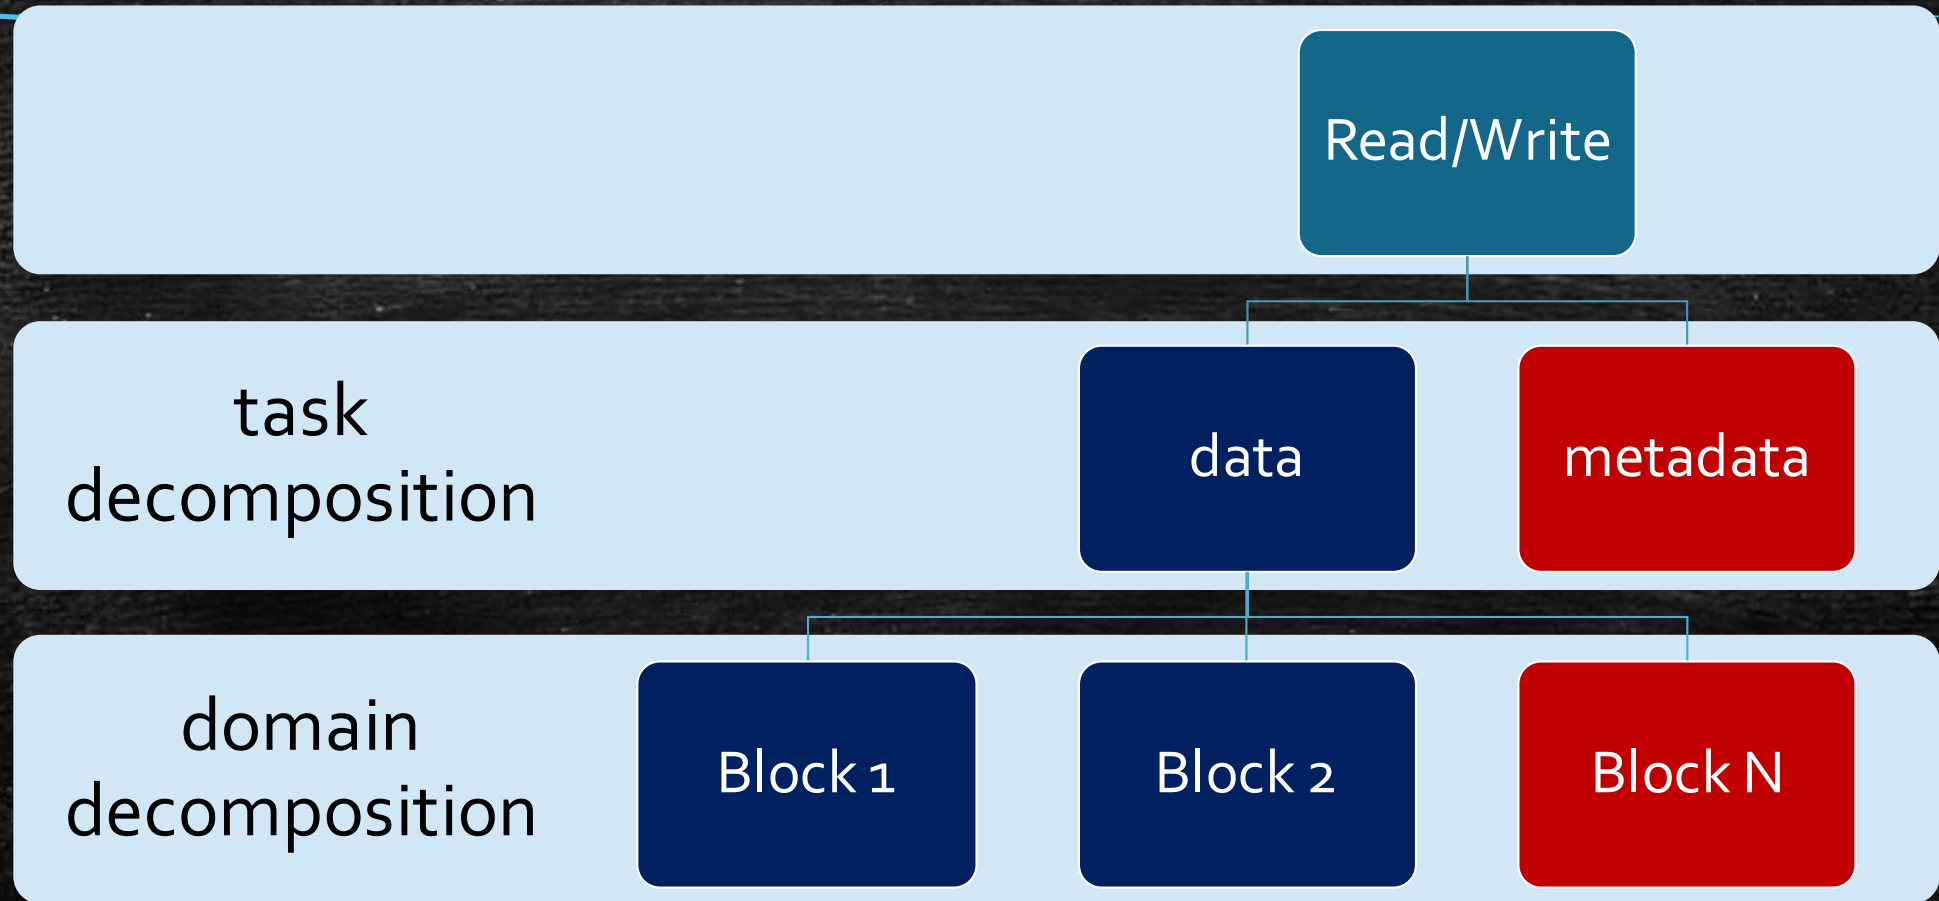
Task Parallelism

Separation of metadata and data operations



Domain Decomposition

Break data into pieces that are to be handled on different hardware



Common Parallel Filesystems



- General Parallel File System (GPFS) / Lustre
 - Commonly associated with supercomputing
 - Employed at NCI, Pawsey, CSIRO and Swinbourne
 - Can be used on AWS and Azure



- Gluster, BeeGFS
 - Parallel distributed file systems
 - More associated with cloud computing

l.u.s.t.r.e.
File System



What is Lustre?

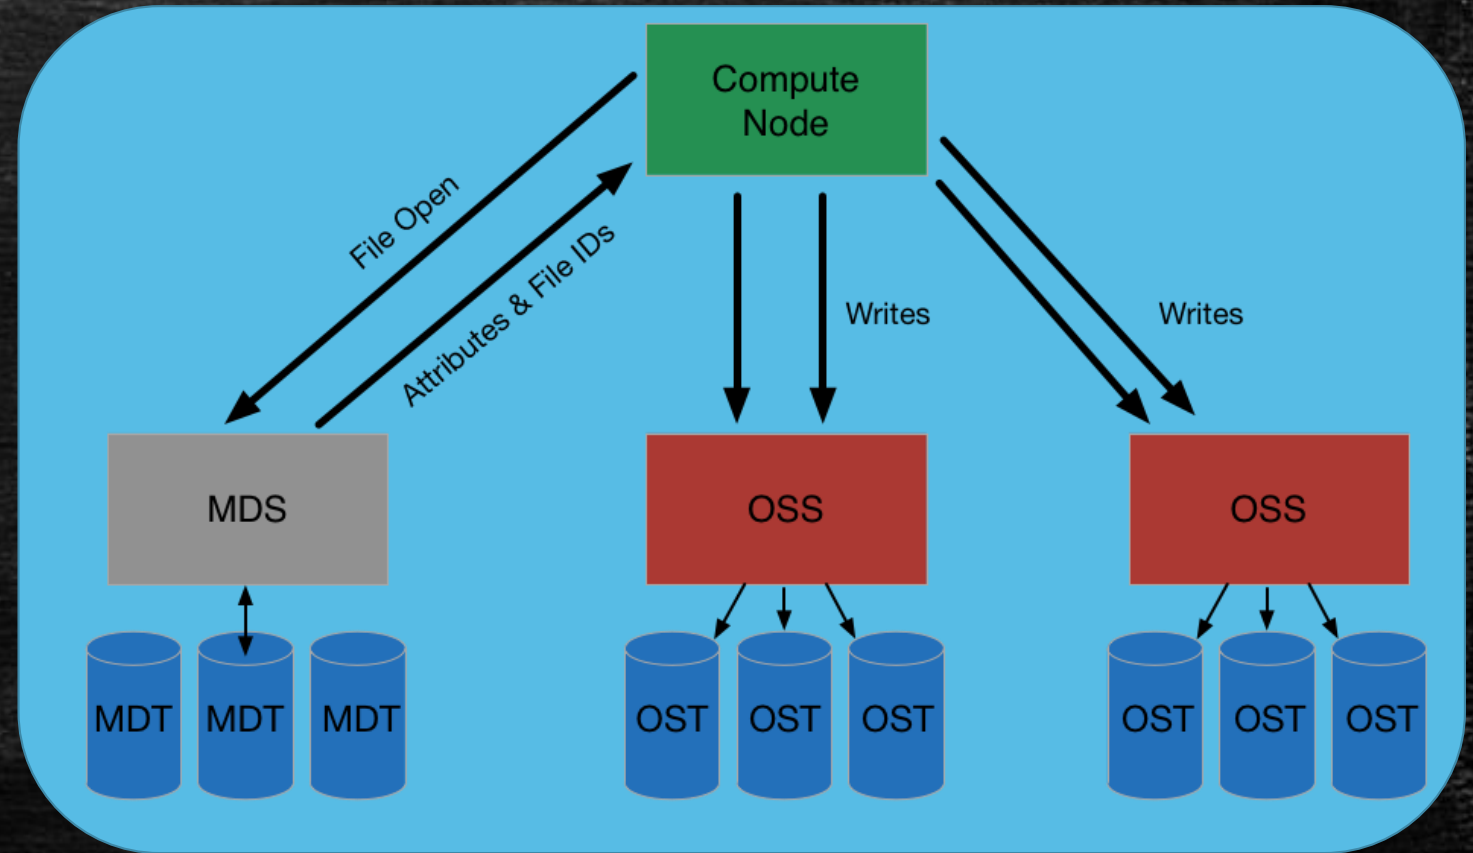
Points of Interest

- Widely used
 - Most hardware vendors offer support
 - Majority of Top 500 systems use it
- Built for heavy duty performance
 - Used with hundreds of thousands of compute nodes
 - Deployed on filesystems with petabytes of capacity
 - capable of Terabyte/sec throughputs
- Industry & community acceptance
 - Lustre users group
 - Application support in libraries such as HDF5
 - Can get service contracts through vendors

Overview

Lustre uses both task and domain parallelism

Different hardware is used for metadata and data operations.



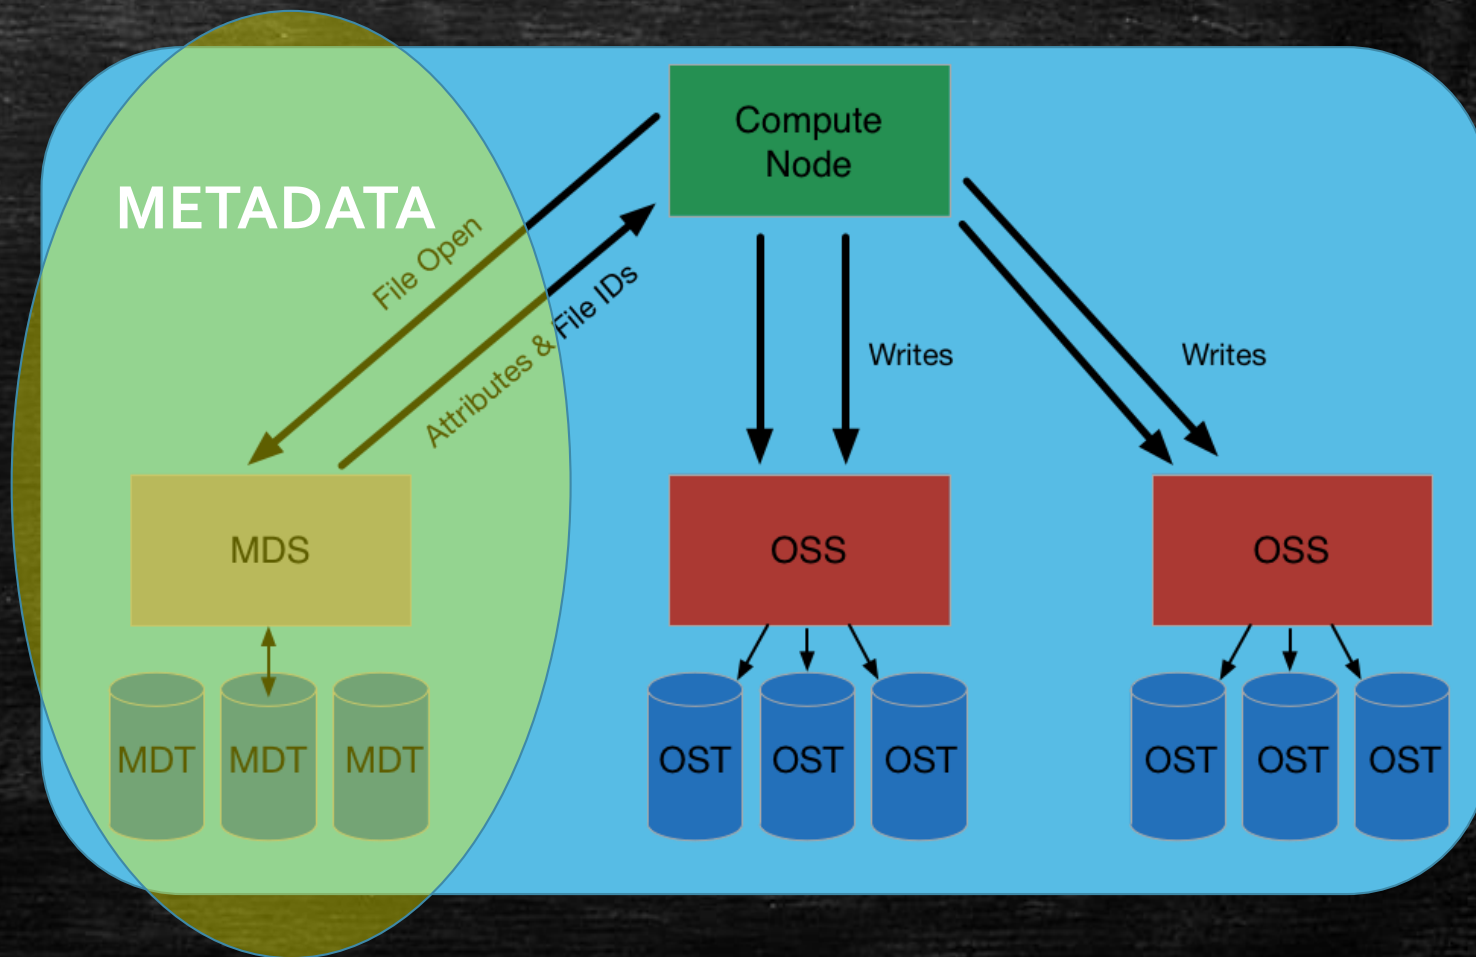
Overview

Meta-Data Server (MDS)

- Tracks storage locations of all files
- Directs I/O requests to appropriate OSTs and OSSes

Meta-Data Target (MDT)

- Stores all metadata
- Filenames, permissions, paths, layout, etc



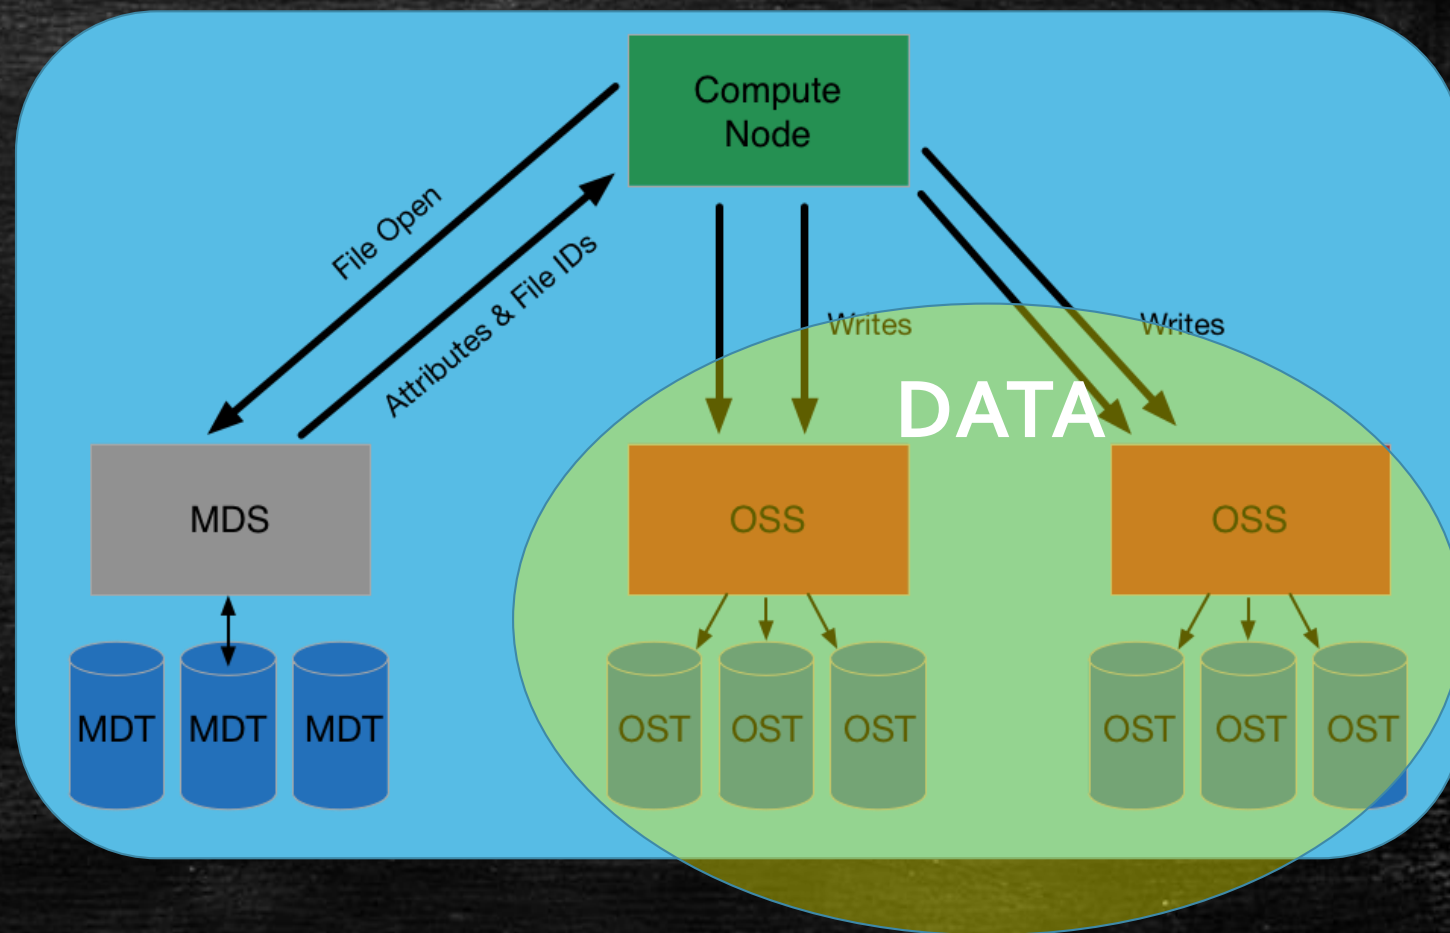
Overview

Object Storage Server (OSS)

- manages a subset of OSTs
- handles network and I/O requests

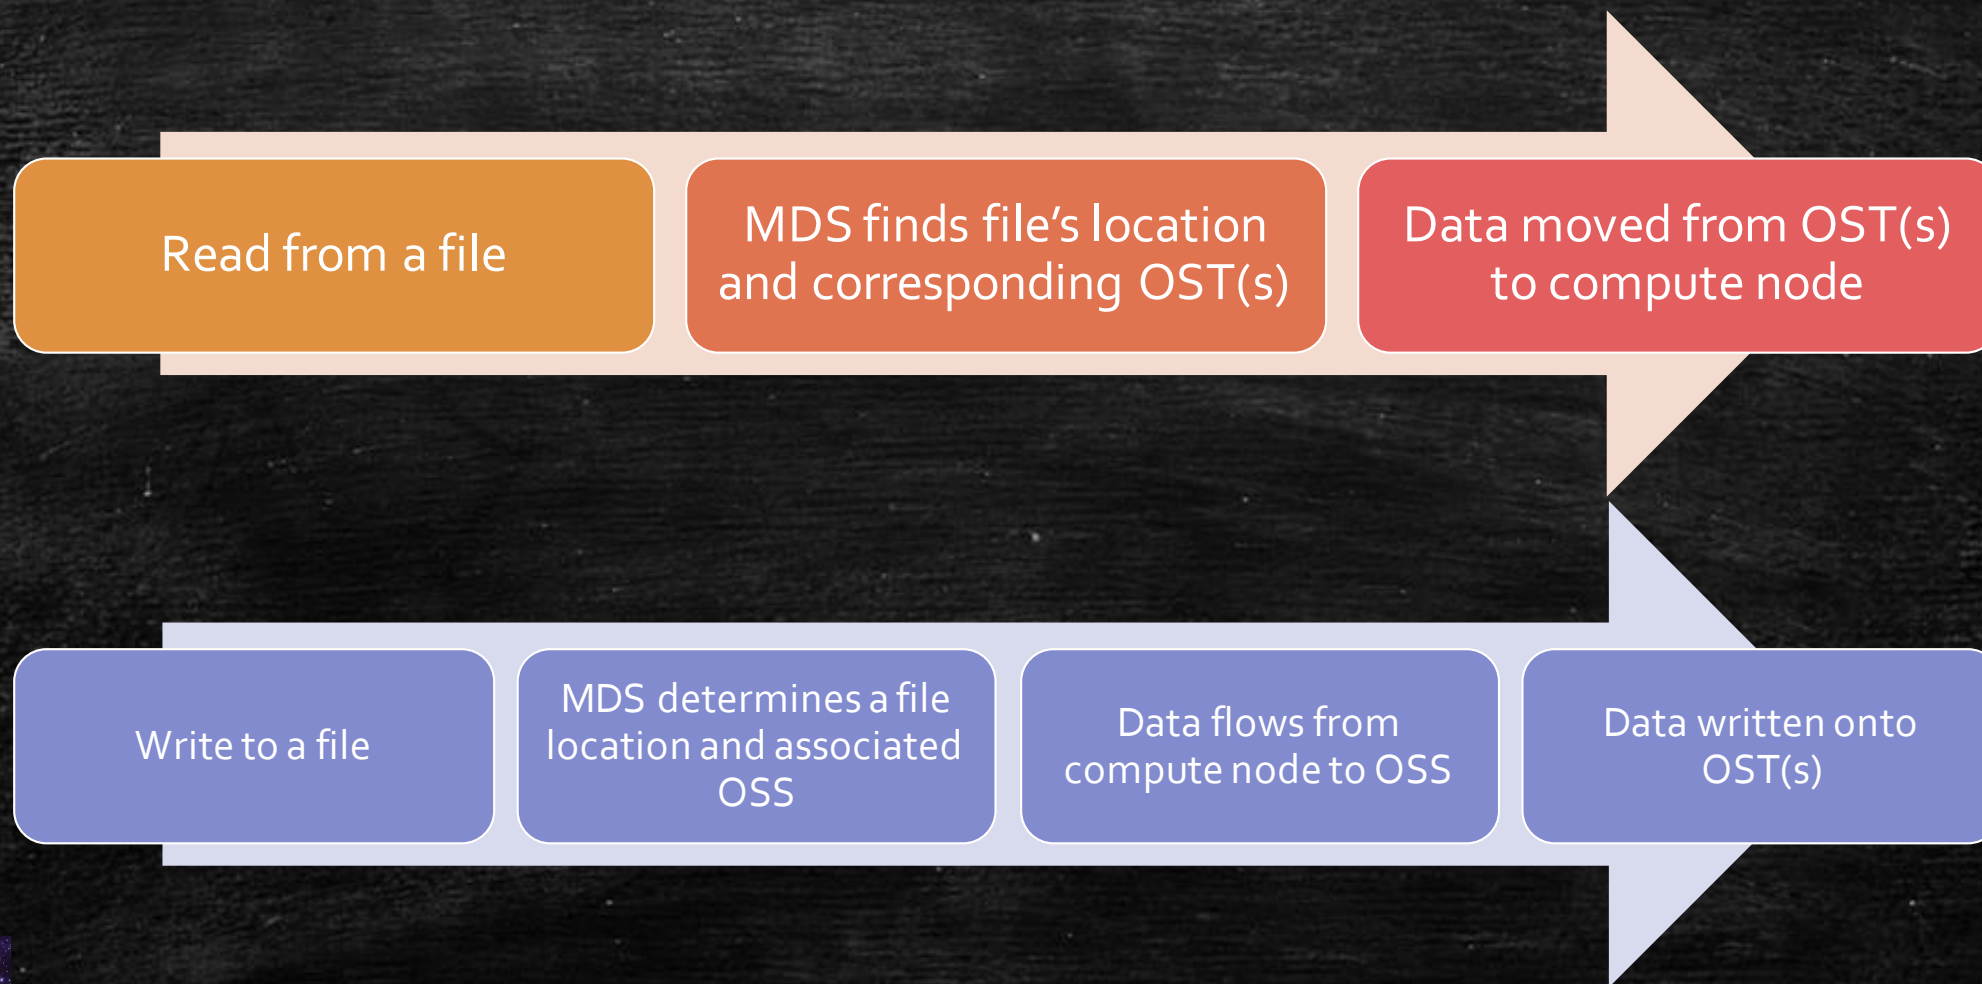
Object Storage Target (OST)

- stores data stripes
- RAID array of hard disks



More hardware dedicated to the DATA side

What happens during a read/write?



Tuneable Parameters in Lustre


Will stick to two parameters from that effect IO performance:




Block Size

The size of the individual data chunks that are routed to/from OSTs.

Too small: there will be too many blocks leading to bandwidth issues



Too large: partially filled blocks are used leading to latency issues



Number of OSTs

A file or directory can have a specified number of OSTs assigned to it. It will have its data blocks spread among these OSTs.

Too few OSTs: Bandwidth issues arise as the data blocks are blocked by the lack of “connections” to hard disk.



Too many OSTs: Latency can become an issue as data blocks are too spread. Metadata tracking can become prohibitively expensive.

Changing Lustre Parameters

- These parameters can be attributed to a single file or a directory of individual files.
- ***Be careful when modifying these parameters. Know what you are doing!***
- System defaults are generally good. Only change when observed IO performance is bad.
- Remember that filesystems are a shared resource!

Demo: Striping and IO performance

Two Python demo scripts are available:

`lustre_block_size_test.py`

`lustre_num_stripes_test.py`

We will investigate the effect that stripe size and number of OSTs play in write performance on a Lustre filesystem.

Which factor has the greatest effect? Why?

Can you think of situations when the each factor is dominant in determining observed IO performance?

Demo: How to Run

1. Click on the terminal in your JupyterHub screen
2. Connect via ssh to the Athena cluster located at Pawsey
`ssh couXXX@athena.pawsey.org.au`
3. Go into the Lustre demo directory
`cd HPC-Workshop/lustre`
4. Submit the 2 demo jobscripts to SLURM
`sbatch jobscript.slurm_block_size`
`sbatch jobscript.slurm_num_osts`
5. Look at the output with the vi text editor
`vi LustreTest_VaryingBlockSize`
`vi LustreTest_VaryingNumOSTs`
6. Exit vi by pressing `:q` and then ENTER

Getting the most out of Lustre

The Good and Bad of Lustre



It excels at:

- large contiguous or uniformly positioned reads/writes
- low frequency IO patterns
- using MPI-IO



It struggles with:

- high frequency reads/writes
- discontinuous / non-uniform read or write patterns
- accessing large numbers of files simultaneously
- Metadata heavy operations (ls -la)

Useful Lustre-specific Commands

Command	Description
<code>lfs mkdir -c <# ost> -s <stripe size> <dir></code>	Create a directory over a given # of OSTs and stripe size
<code>lfs getstripe <file / dir></code>	Get the stripe size and # of OSTs assigned to a file or directory
<code>lfs setstripe -c <# ost> -s <stripe size> <file></code>	Set the stripe size. # is number of bytes
<code>lfs migrate -c <# ost> -s <stripe size> <file></code>	Modify the stripe count and/or size for an existing file or directory
<code>lfs find <dir></code>	List the contents of a directory
<code>lfs df</code>	Check disk space usage

Other recommendations

- **Avoid** doing the following:
 - using auto/tab complete
 - wildcard expressions (especially `tar` and `rm`)
 - Generate an explicit list of files and operate on each individually
 - thousands of files in a single directory
 - use a directory structure
 - Using `ls -l` unless absolutely necessary
- Open files as `READ_ONLY` if they will only be accessed for reading
- Don't stripe small files (<1 MB) across multiple OSTs
 - Set stripe count to 1 for a directory and write files in it
 - Be careful not to specify a specific OST!!

Other recommendations

- Store executables in non-parallel file system
- Don't edit text/source files in a Lustre file system
 - `vi` creates temporary "snapshot" files that increase metadata loads
- Remember to appropriately stripe when moving files from a non-parallel filesystem to the Lustre filesystem
- System-level caching can effect expected performance