



# Programming Paradigms

---

Mark Cheeseman  
*Mark.Cheeseman@csiro.au*



# Overview



## Types of parallelism

- task parallelism
- domain decomposition



## Memory models

- shared vs distributed
- state
- *hybrid\**

# Parallelism

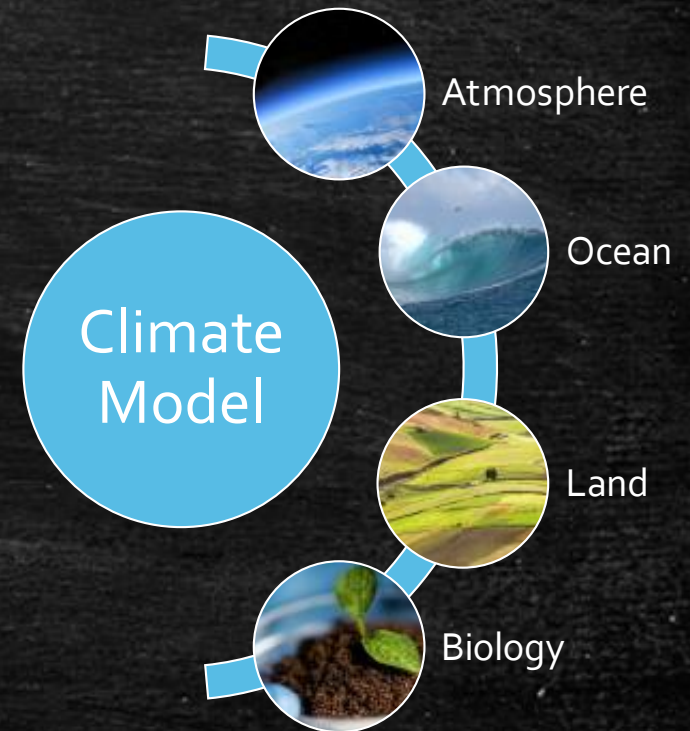
---





# Task Parallelism

- Occurs where clearly distinguishable operations can be run independently.
  - Input/output of each process may depend on another process BUT there is no/little data dependence during the operation.
- Can be in a single application or a workflow
- Example would be a coupled climate model.
  - Independent binaries for ocean, atmosphere, ice, land and biological models are run with minimal interaction.





# I/O Server Design

---

Assume we have an application running concurrent tasks that each have data we wish to write to hard disk. We can

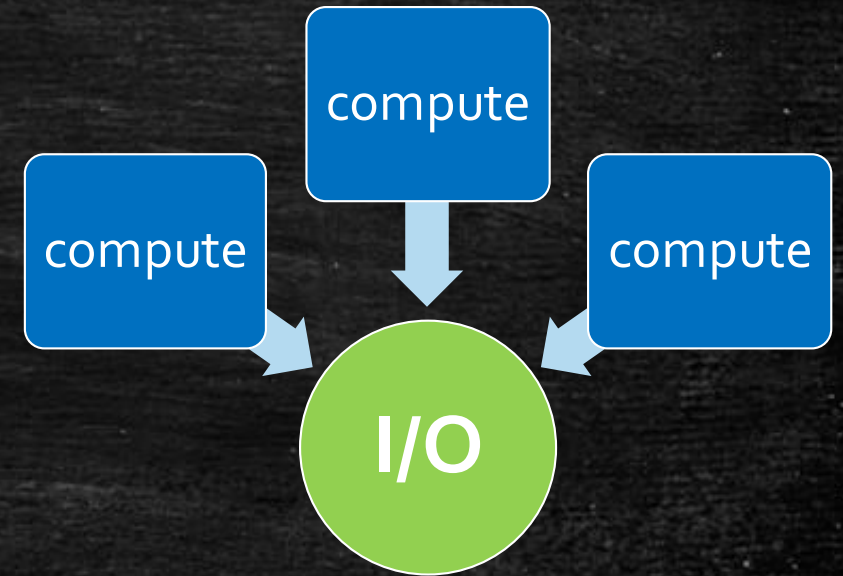
- a) have each task write its own output file
  - messy, unmanageable as # of tasks grows
- b) Assign another concurrent task to be responsible for the collection and writing of output data from all tasks (*io server approach*)
- c) Use a parallel-IO enabled infrastructure
  - Learn more on Day 3!



# I/O Server Design

---

- Allows a single output file to be written
- Implementation details are key!
  - Efficient transferring of data from compute to I/O
  - Will I/O cope with the workload?
  - How would we grow this design to deal with
    - More compute tasks?
    - Larger output data requests?
    - More frequent data requests?



# Domain Decomposition

---

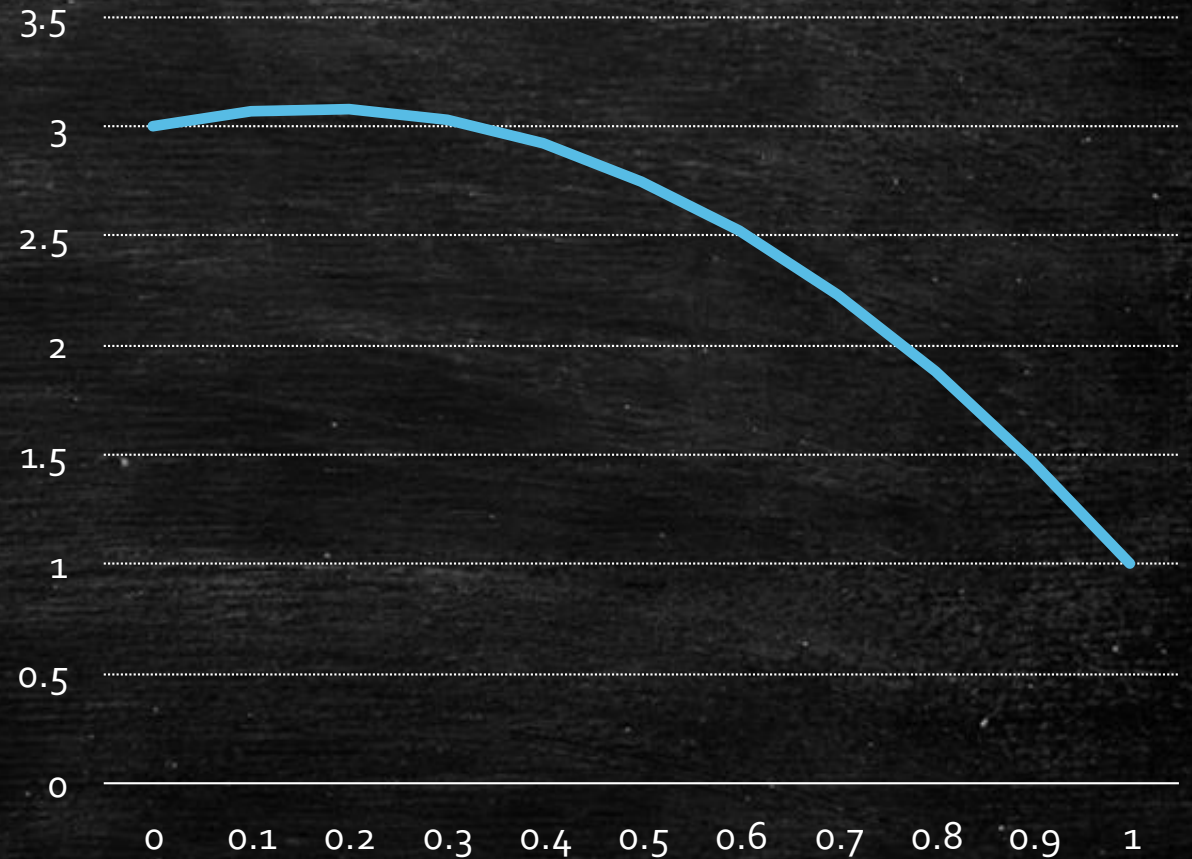
- Want to perform an operation on a large dataset.
- Identical instances of this operation can be performed independently
  - Data may need to be transferred between the instances
- Most commonly used form of parallelism in software
- Examples:
  - vector operations
  - Finite difference/element codes



# Numerical Integration

- Assume we have smooth continuous function over a finite interval  $[0, 1]$
- Want to evaluate its integral over the interval.

$$\int_0^1 f(x) dx$$

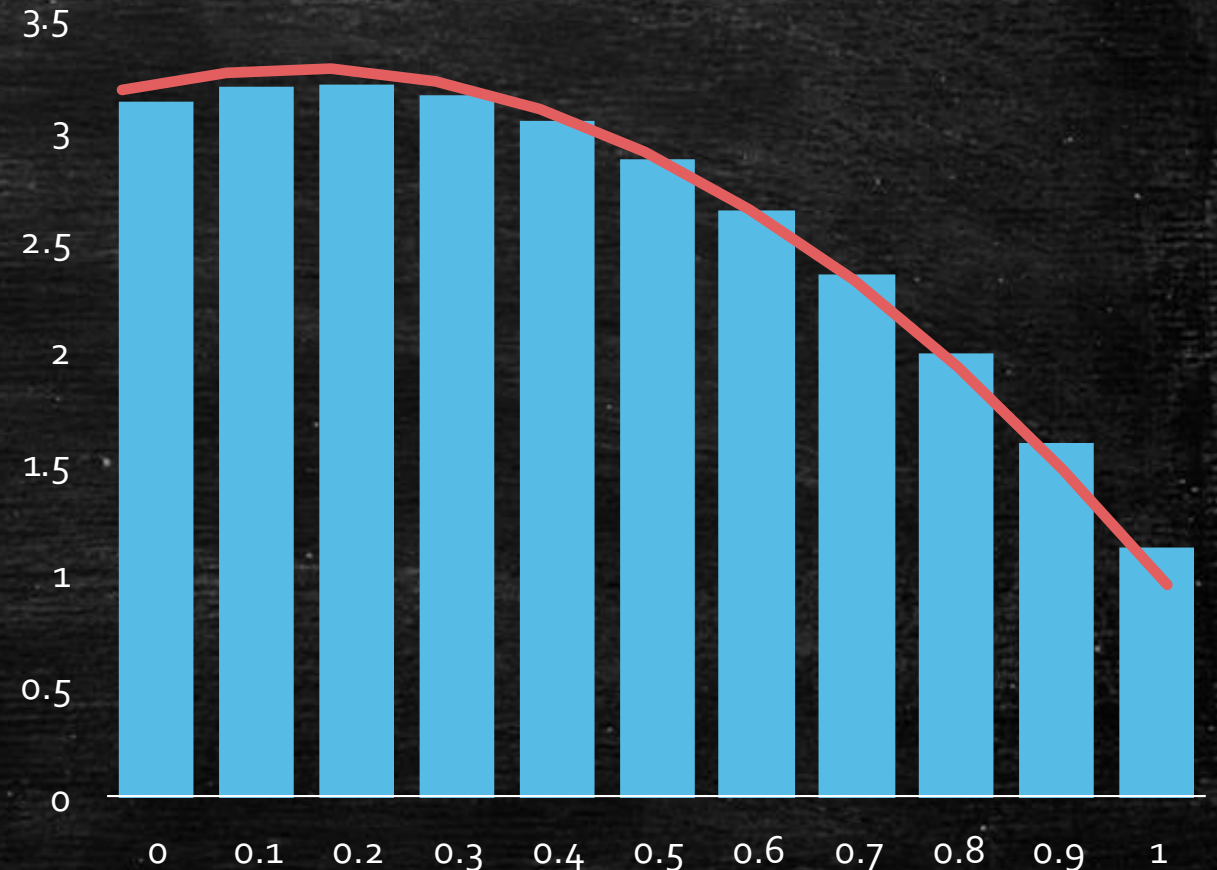




# Numerical Integration

- Integral can be approximated by the area under the curve.
- Can determine area by summing up a series of rectangles.

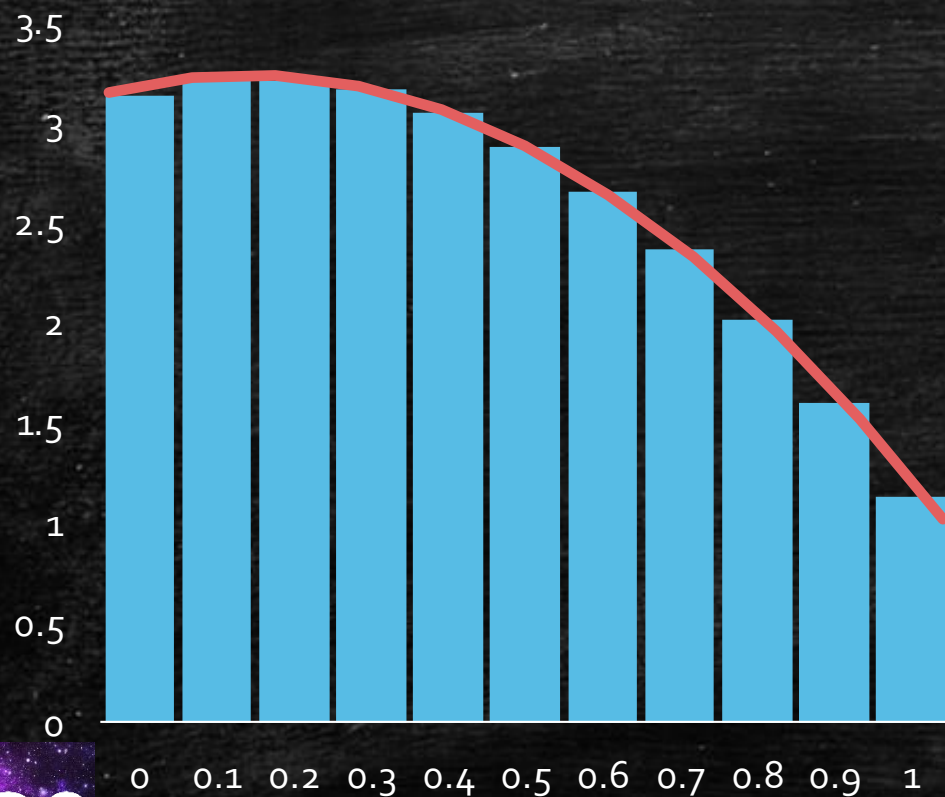
$$\int_0^1 f(x) dx \approx \sum_{i=1}^N a_i$$



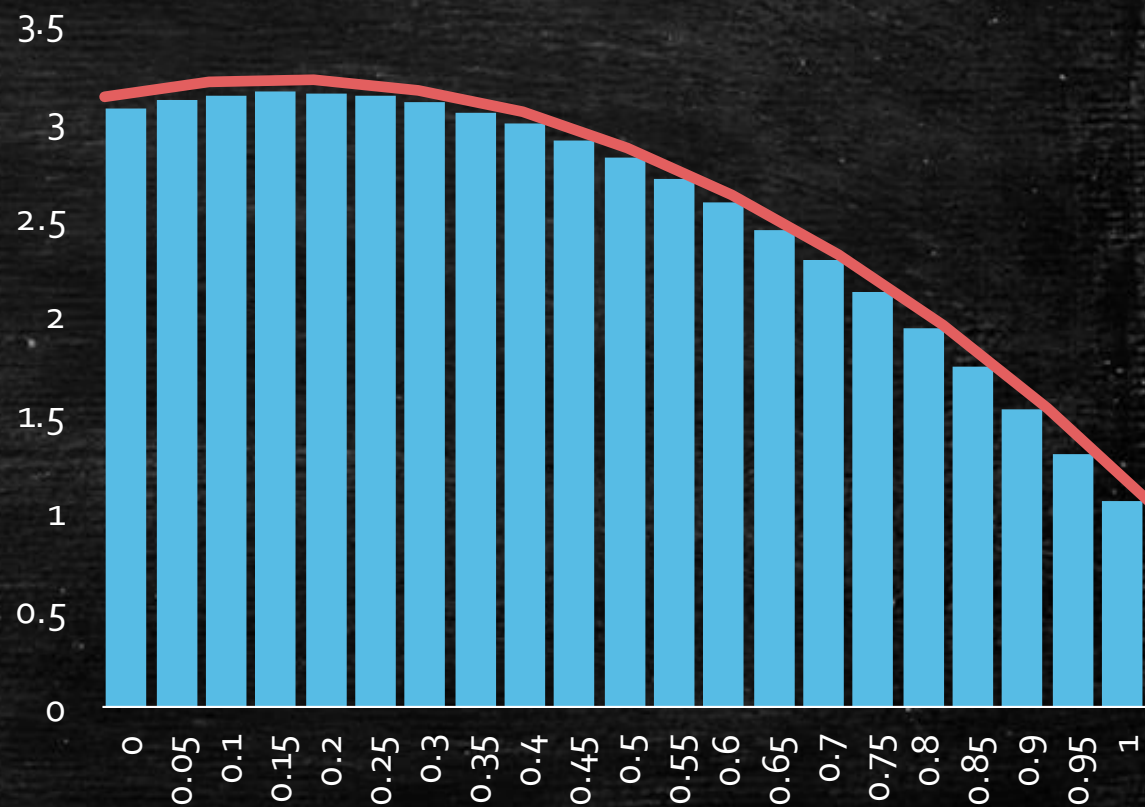


# Numerical Integration

- Solution accuracy increases by using more finer rectangles



- More rectangles leads to higher computational workload

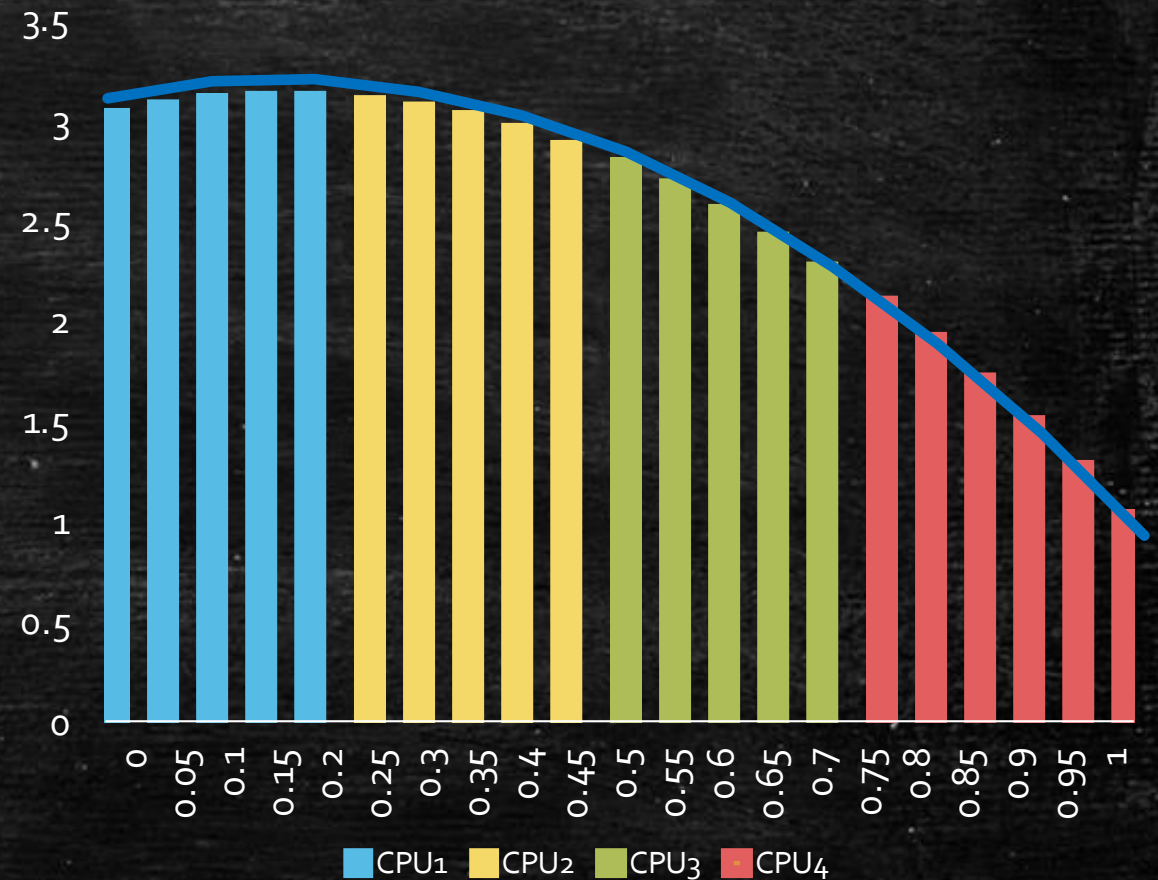




# Numerical Integration

- Assign each process to compute area for a piece of the global domain
- Add up the individual areas to approximate the global integral

$$\int_0^1 f(x) dx \approx A1 + A2 + A3 + A4$$





# Numerical Integration

---

1. Define/read in  $f(x)$  and interval of interest  $[x_1, x_2]$
2. Define sub-intervals for each concurrent task
3. Every task determines the sum of areas for all rectangles it owns
4. All sums are gathered onto a single (or all) tasks



# Numerical Integration

---

1. Define/read in  $f(x)$  and interval of interest  $[x_1, x_2]$
2. Define sub-intervals for each concurrent task
3. Every task determines the sum of areas for all rectangles it owns
4. All sums are gathered onto a single (or all) tasks

Can be expensive and difficult



# Numerical Integration

---

1. Define/read in  $f(x)$  and interval of interest  $[x_1, x_2]$
2. Define sub-intervals for each concurrent task
3. Every task determines the sum of areas for all rectangles it owns
4. All sums are gathered onto a single (or all) tasks

Load balancing:

- what if # of tasks doesn't evenly divide into interval?
- what if CPUs are not homogeneous?



# Why are task and domain parallelism important?

---

Many codes/workflows tend to naturally fit one. Factors affecting your choice include:

- Domain size
- Task independence / complexity
- Machine architecture (lots of cores tends towards data parallelism)
- Software used (programming language, libraries)
- Your objectives (ease of use, run time, memory usage, etc)



# Memory Models

---





# Memory Models

---

Shared



Distributed



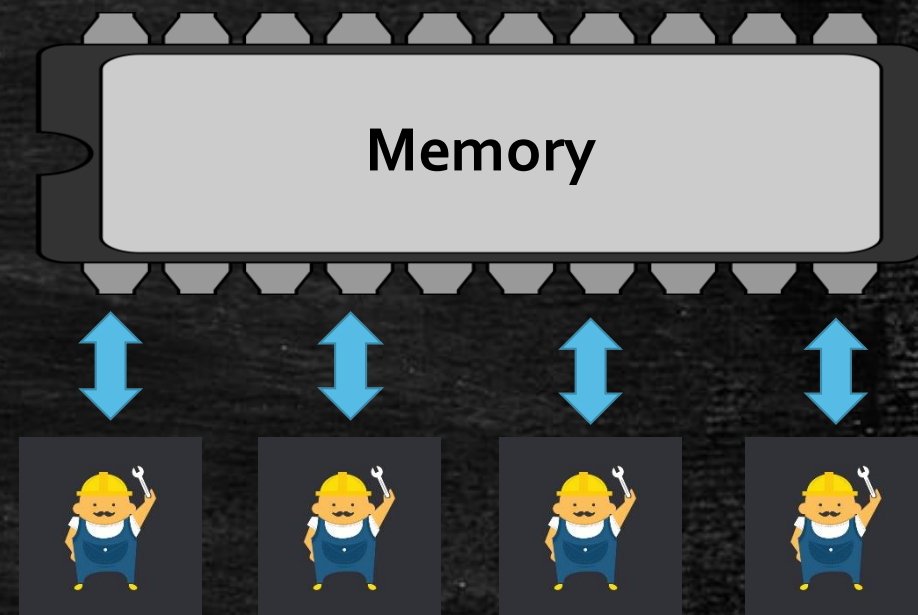
*Hybrid*





# Shared memory model

- Description
  - multiple processes that have access to a single global memory space
  - *everyone sees everything*
- Implementations
  - OpenMP
  - Threaded programming (Java)
  - CUDA





# Shared memory model

---



## Cons

- state protection.
- problem size limited to size of memory available
- number of processes limited by platform
- performance issues may be harder to address
  - memory layout and/or hierarchies
  - architecture features



## Pros

- simple to program
  - Extra coding kept to a minimum
  - Native compiler/package support
- fast
  - Everything done on local platform (CPU, GPU, socket)
- generally available



# Memory and State

---

- All processes can see and access any variable in the shared memory

```
int tid, x = 1;

#pragma omp parallel {

    tid = omp_thread_num() + 1;

    x += tid;

    printf( "value of x is %f\n", x );

}
```

What is the printed value of x? (assuming 4 threads)

What would the value of x be outside of the parallel region?



# Memory and State

---

- Inside the parallel region:
  - Don't know the value of  $x$  at a particular time
  - Value of  $x$  depends on the thread scheduling (which is compiler and run specific)
- Outside the parallel region:
  - $X = 11$
  - Addition is an associative operation so the thread scheduling has no effect on the final outcome.
  - What could happen if we were adding a large series of very small and very big numbers?
    - *Remember that we are working with a finite number system!*



# Memory and State

---

- Need to explicitly monitor access to shared variables
  - Otherwise unintended consequences could occur!
- Be careful with non-associative operations
  - Order of operations would be important!



# Numerical Integration – Shared Memory

---

1. Define/read in  $f(x)$  and interval of interest  $[x_1, x_2]$
2. Define sub-intervals for each concurrent task
  - a. *Spawn 4 threads*
  - b. *Each thread would determine its sub-interval using its thread ID and global interval limits*
3. Every task determines the sum of areas for all rectangles it owns
4. All sums are gathered onto a single (or all) tasks



# Numerical Integration – Shared Memory

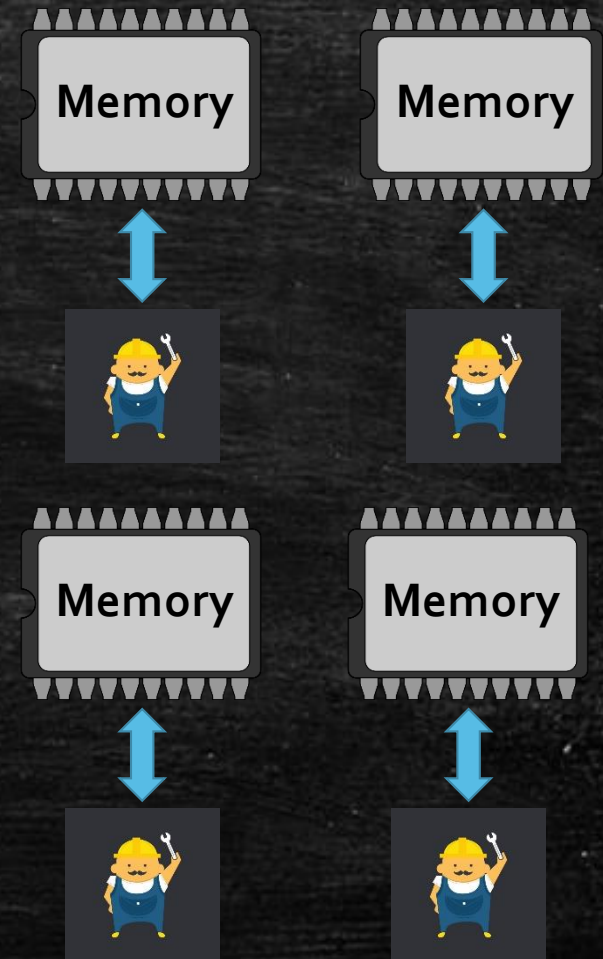
---

1. Define/read in  $f(x)$  and interval of interest  $[x_1, x_2]$
2. Define sub-intervals for each concurrent task
  - a. *Spawn 4 threads*
  - b. *Each thread would determine its sub-interval using its thread ID and global interval limits*
3. Every *thread* determines the sum of areas for all rectangles it owns
4. All sums are gathered onto a single (or all) tasks
  - a. *Performed as a local addition. Order is irrelevant.*



# Distributed memory model

- Description
  - multiple processes located on physically separated platforms each with their own memory space. Global memory space is not physically contiguous.
  - *each process only sees its part of the problem*
- Implementation
  - Message Passing Interface (MPI)
  - Partitioned Global Address Space (PGAS) languages
    - Co-array Fortran, Charm++, Global Array Toolkit





# Memory and State

---

- Each individual compute node has its own copy of variables
- Be careful about local and expected global values



# Numerical Integration – Distributed Memory

---

1. Define/read in  $f(x)$  and interval of interest  $[x_1, x_2]$ 
  - a. *Start up the distributed environment with 4 active tasks*
  - b.  *$f(X)$  and defined on every task – OR – defined on a root task and then distributed*
2. Define sub-intervals for each concurrent task
3. Every task determines the sum of areas for all rectangles it owns
4. All sums are gathered onto a single (or all) tasks



# Numerical Integration – Distributed Memory

---

1. Define/read in  $f(x)$  and interval of interest  $[x_1, x_2]$ 
  - a. *Start up the distributed environment with 4 active tasks*
  - b.  *$f(X)$  and defined on every task – OR – defined on a root task and then distributed*
2. Define sub-intervals for each concurrent task
3. Every task determines the sum of areas for all rectangles it owns
4. All sums are gathered onto a single (or all) tasks
  - a) *Need communication between ALL tasks*
  - b) *Higher overhead than with the shared memory approach*



# Distributed memory model

---



## Cons

- more complex code
  - elaborate API to follow. Harder to visualize.
- slower passing of information between processes
  - Order of magnitude slower in shared memory case



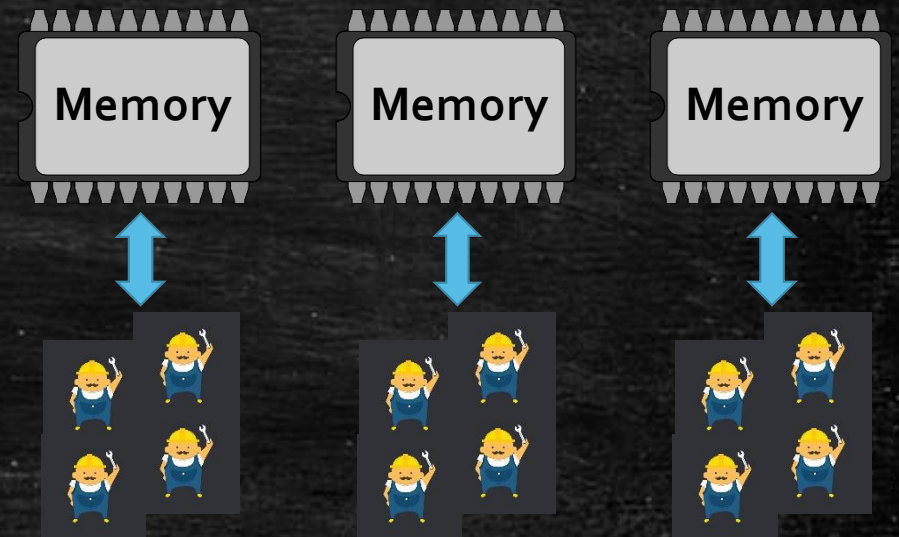
## Pros

- State problem is easier
- acquire more resource when needed
  - more CPUs for smaller run times
  - more memory for larger domains
- Universal acceptance
  - MPI found on almost every platform
  - lots of libraries, tools and packages
- Other features
  - Parallel I/O
  - Multi-platform support (CPU + GPU)



# Hybrid - Multi-core

- As # of cores per CPU grows, combining task and data parallelism is required.
- Ex: adding OpenMP loop unrolling to MPI tasks
- Most commonly used to overcome bandwidth issues.
  - network saturation between CPUs
  - I/O





# Hybrid - GPU Computing

- GPUs and other accelerators (Xeon Phi) use a hybrid parallelism as well
  - Host and GPU device have separate memories
  - Both are capable of running multiple processes
- Modern GPGPUs are
  - Great for data parallelism (lots of threads)
  - Good for task parallelism as well (use of queues)
- Memory issues are paramount

**OpenACC**  
Directives for Accelerators

